

Regularization strategies for Deep Learning

Weight Decay, Augmentation, Early Stopping,
Dropout & more


N. Rich Nguyen, PhD
SYS 6016



**THE BEST WAY TO
EXPLAIN OVERFITTING**

Review: Basic Concepts

- The ability to perform well on *unobserved inputs* is called **generalization**.
- **Underfitting**: when the model is not able to obtain sufficient training error.
- **Overfitting**: when the gap between training and generalization error too large.
- **Bias** and **Variance** Tradeoff: per your discussion on Piazza.




Review: Model Capacity


A model capacity is its ability to fit a wide variety of functions:

- Models with **low capacity** may struggle to fit the training set (Underfitting)
- Models with **high capacity** can overfit by memorizing properties of the training set that do not serve them well on the test set (Overfitting)


One way to control model capacity is by choosing its **hypothesis space**, the set of functions that the learning algorithms is allowed to select as being the solution.



$$\hat{y} = \theta_0 + \theta_1 x_0$$



$$\hat{y} = \theta_0 + \theta_1 x_0 + \theta_2 x_0^2$$




$$\hat{y} = \theta_0 + \sum_{i=1}^9 \theta_i x_i$$

Regularization

Regularization is the modification made to a learning algorithm that is intended to reduce its generalization error but not its training error.

For example, we can modify the loss function for linear regression to include a preference for the weights to have smaller L2 norm as a **regularizer**:

$$J(\mathbf{w}) = \mathcal{L}_{\text{mse}}(\mathbf{X}_{\text{train}}, \mathbf{y}_{\text{train}}) + \lambda \mathbf{w}^T \mathbf{w}$$



Overview: Regularization strategies

We will survey the general strategies used to regularize neural networks

1. Parameter Norms
2. Dataset Augmentation
3. Multi-task Learning
4. Early Stopping
5. Bagging
6. Dropout
7. Tangent Prop

1. Parameter Norms

Parameter Norm Penalties

Many regularization approaches are based on **limiting the model capacity** by adding a **parameter norm** penalty to the objective (loss) function:

$$J(\mathbf{w}; \mathbf{X}, \mathbf{y}) = \underbrace{\mathcal{L}(\hat{\mathbf{y}}, \mathbf{y})}_{\text{Data Loss}} + \underbrace{\lambda \Omega(\mathbf{w})}_{\text{Norm Penalty}}$$

where λ is a hyperparameter that controls the relative contribution of the norm penalty term, Ω , relative to the standard data loss function \mathcal{L}

- Larger value of λ correspond to more regularization
- Setting λ to 0 results in no regularization
- Norm penalty Ω penalizes only the weights of the affine transformation
- Different choice of Ω can result in different solutions being preferred.

L² Parameter Regularization

L² regularization is aka **Ridge Regression** or **Tikhonov regularization**

The L² **norm** penalty commonly known as **weight decay**

$$J(\mathbf{w}; \mathbf{X}, \mathbf{y}) = \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y}) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

$$\nabla_{\mathbf{w}} J(\mathbf{w}; \mathbf{X}, \mathbf{y}) = \nabla_{\mathbf{w}} \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y}) + \lambda \mathbf{w}$$

To take a single **Gradient Descent** step to update the weights:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \nabla_{\mathbf{w}} J(\mathbf{w}; \mathbf{X}, \mathbf{y})$$

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha (\nabla_{\mathbf{w}} \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y}) + \lambda \mathbf{w})$$

$$\mathbf{w} \leftarrow (1 - \alpha\lambda) \mathbf{w} - \alpha \nabla_{\mathbf{w}} \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y})$$

< 1

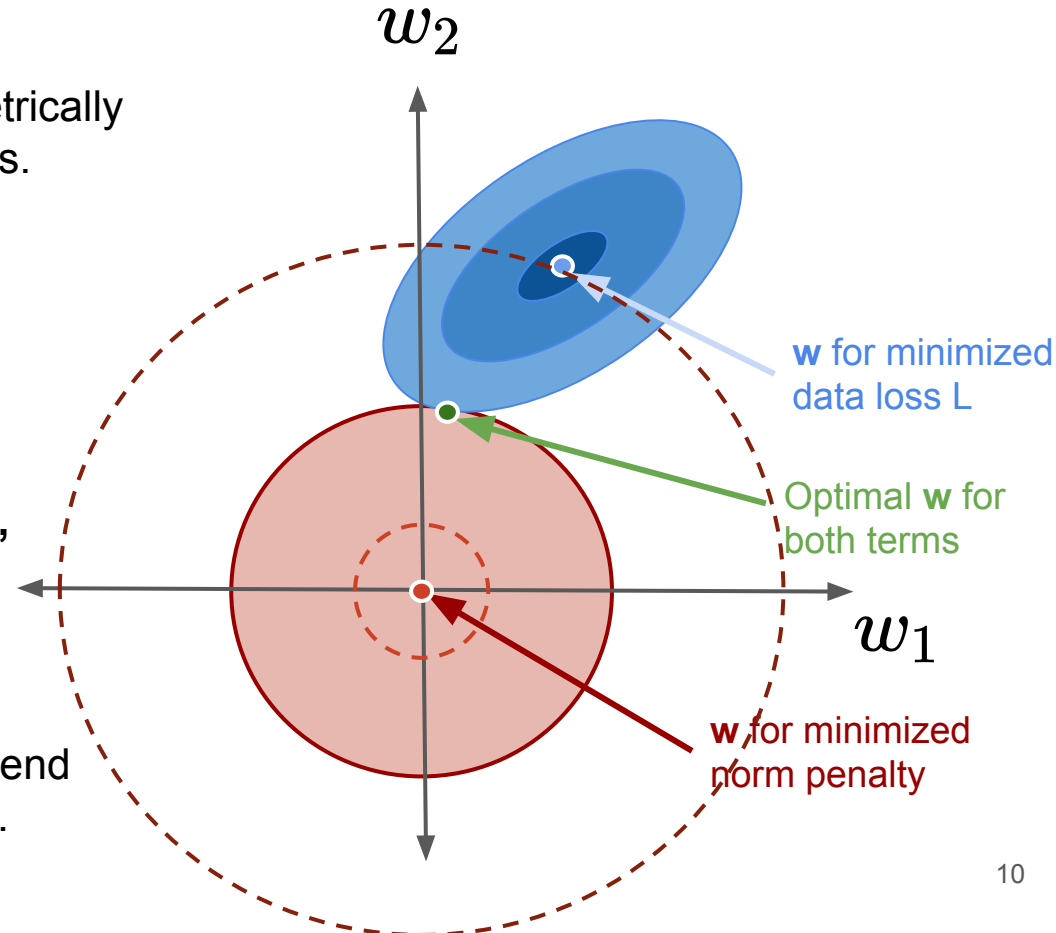
Shrink the weight vector before gradient update

**HONEY, I
SHRUNK
THE KIDS
COLLECTION**



Geometry Interpretation

- L^2 regularizations $||\mathbf{w}||^2$ can be geometrically represented as concentric (red) circles.
- When circle is too small, the params are not useful to the model.
- When the circle region (in red) grows, due to its shape the region intersects the data contour **closer to the origin**, L^2 makes both parameters shrink and w_1 near zero.
- When the circle grows too large, you end up with a similar params as data loss.



L¹ Parameter Regularization

L¹ regularization is aka **LASSO** (least absolute shrinkage and selection operator)

L¹ norm commonly is known as the **Manhattan Distance**

$$J(\mathbf{w}; \mathbf{X}, \mathbf{y}) = \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y}) + \lambda ||\mathbf{w}||_1$$

$$\nabla_{\mathbf{w}} J(\mathbf{w}; \mathbf{X}, \mathbf{y}) = \nabla_{\mathbf{w}} \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y}) + \lambda \text{sign}(\mathbf{w})$$

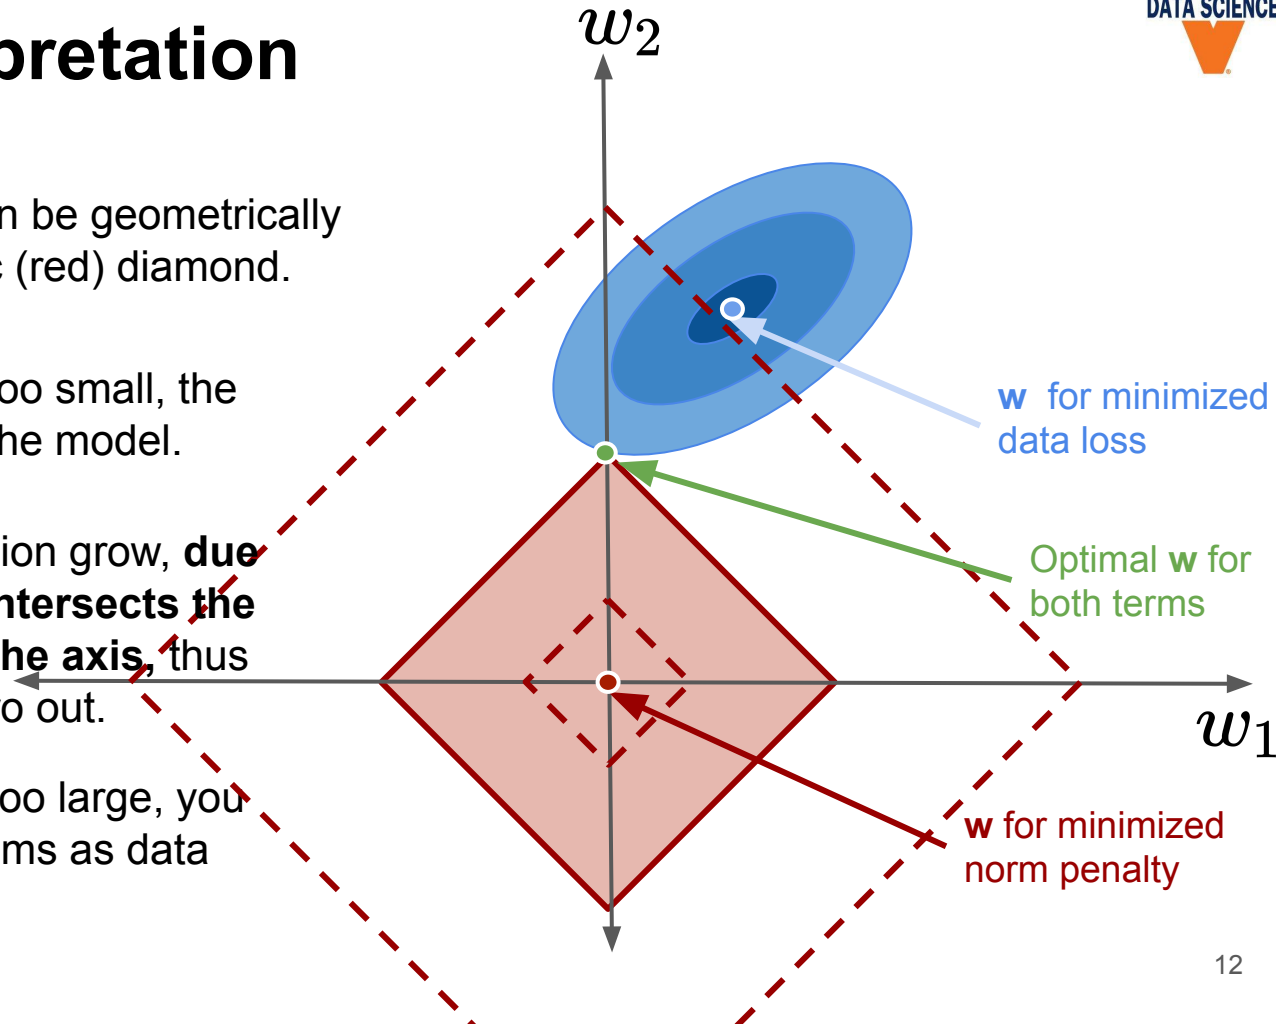
The regularization to the gradient no longer scale linearly with each w , instead it is a constant factor with a sign equal to $\text{sign}(w)$. Thus, there is **no clean** algebraic solution to approximates J as we have just seen in L² regularization

L¹ norm makes the parameters to become **sparse** (contains lots of zeros)

L¹ norm tends to cause a subset of the network weights to become zero, suggesting that the corresponding signal may safely be discarded (**dead neuron**).

Geometry Interpretation

- L^1 regularizations $\|\mathbf{w}\|_1$ can be geometrically represented as concentric (red) diamond.
- When diamond region is too small, the params are not useful to the model.
- When as the diamond region grow, **due to its shape the region intersects the data contour on one of the axis**, thus the other parameter is zero out.
- When diamond region is too large, you end up with a similar params as data loss.



Can we combine multiple regularizations?


Elastic Net

Is a middle ground between Ridge (L2) and Lasso (L1) with a mix ratio r

$$J(\mathbf{w}; \mathbf{X}, \mathbf{y}) = \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y}) + \underbrace{r\lambda ||\mathbf{w}||_1}_{\text{LASSO}} + \underbrace{(1-r)\lambda ||\mathbf{w}||_2^2}_{\text{RIDGE}}$$


- Preferable to have at least a little bit of regularization, **Ridge** if a good default
- If you suspect that only a few features are actually useful, use **Lasso**
- Lasso may behave erratically when `#features > #examples`, use **Elastic Net**

Geometry Interpretation



Family of Parameter Norm Models


$$J(\mathbf{w}; \mathbf{X}, \mathbf{y}) = \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y}) + \lambda ||\mathbf{w}||_q$$



2. Dataset Augmentation

Dataset Augmentation

- Make the model generalize better by training it on **more data**
- If the amount of data available is limited, get around by creating “fake” data and augment them into the training set.
How?
- Augmentation has been particularly effective for object recognition. Image transformations (ie. translating, rotating, cropping, brightness correcting, ect) often greatly improve generalization.



Injecting noise and label smoothing

- Dataset augmentation is **not** as readily applicable for many other tasks. For example: apply transformation that would change to correct class (ie. 180 rotations are not appropriate for “b” and “d”; or “6” and “9”)
- **Injecting (random) noise** in the inputs can be a form of data augmentation.
- Injecting noise to output labels? It can be harmful to maximize $\log p(\underline{y} | \mathbf{x})$ when \underline{y} is a mistake.
- Prevent this by explicitly model the noise on the label: \underline{y} is correct with probability $1 - \epsilon$ (epsilon) with small constant ϵ
- **Label smoothing** regularizes a model on a softmax with k output values by replacing hard label 0 and 1 with $\frac{\epsilon}{k-1}$ and $1 - \epsilon$
- Label smoothing prevents the pursuit of hard probabilities without discouraging correct classification, and is shown prominently in modern neural nets