

Optimization for Deep Models

Initialization, Activation, and Normalization

N. Rich Nguyen, PhD
SYS 6016

0. Challenges

Optimization in ML

Optimization: finding the parameters θ of a model that significantly reduce a cost function $\mathcal{J}(\theta)$, which typically includes a performance measure evaluated on the entire training set as well as additional regularization terms.

The simplest way to convert a ML problem back into an optimization problem is to minimize the expected lost on the training set → **empirical risk minimization**

$$\mathbb{E}_{\mathbf{x}, y} [L(f(\mathbf{x}; \theta), y)] = \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}; \theta), y^{(i)})$$

Rather than optimizing the **risk directly**, we optimize the empirical risk and hope that the risk decreases significantly as well.

Optimization for Neural Networks

Unfortunately, the empirical risk minimization is prone to **overfitting** as model with **high capacity** such as a neural net can simply memorize the training set.

In deep learning, we use a slightly different approach here → **mini-batch stochastic methods**




- Mini-batch sizes are driven by hardware, parallel, memory. It is common for power of 2 as batch sizes (32--256)
- It is crucial that the mini-batch is selected **randomly** → **stochastic**

Compute the gradient of the loss w.r.t. the parameters for that minibatch, then updating the parameters in the direction of the gradient:

$$\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), y^{(i)})$$

$$\theta \leftarrow \theta - \epsilon \mathbf{g}$$

Local Minima, Plateaus, and Saddle Points



- With non-convex functions, such as neural nets, it's possible to have many **local minima**
- Local minima with high cost could pose a serious problem for gradient-based algo
- High-D non-convex functions, another point with zero gradient: **a saddle point**
- Some points around a saddle point have **greater** cost while others have **lower** cost

Training Deep Models

To tackle complex problems such as a visual perception task, we usually need to train a much **deeper network** with hundreds of layers and thousands of neurons.

A few issue with deep models includes:

- Training will be extremely slow
- Millions of parameters risking overfitting
- Facing the tricky “vanishing gradients” problem?

The Vanishing Gradient Problem

- Backpropagation works by passing the **error gradient** back and forth among input, output, and hidden layers. It uses gradient to update the parameters
- For **deep** computational graph, repeated applying the same operation at each time step makes the gradients get **smaller and smaller** to which point they leave the weights virtually **unchanged** → **vanishing gradient problem**



$0.5^{100} =$
7.8886091e-31

The Exploding Gradient Problem

- The opposite of vanishing gradient can also happen
- **Repeated** matrix multiplication at each time step of a computational graph
- Gradients grow bigger and bigger → many layers got **insanely large** weight updates, and the network becomes unstable and diverged



$1.5^{100} =$
 $4.0656118e+17$

A reason why deep neural networks were abandoned for 2 decades!

Optimization of Deep Neural Nets

1. **Initialization:** how to initialize the weights so that they do not saturate?
2. **Activation:** how to solve the vanishing gradient problem?
3. **Normalization:** how to get the model to learn the optimal scale?
4. **Optimizers:** when gradient descent was too slow or not good enough?
5. **Adaptive Learning Rate:** what if convergence is too slow or sub-optimal?
6. **Second-Order Training Methods:** can we make use of second derivatives?



1. Initialization



Xavier Initialization

- Iterative algorithms require the user to specify some initial point from which to begin the iterations → deep learning are **strongly** affected by **initialization**
- We don't want the signal to **vanish**, **explode**, nor **saturate** → break **symmetry** between different units.
- To keep it **flow symmetrically**, Xavier Glorot and Yoshua Bengio argue that “we need the variance of the **outputs** of each layer to be **equal** to the variance of its **inputs**” → Xavier Initialization for Logistic Activation (2010)
- This heuristic strategy led to the current success of Deep Learning.

$$\mathbf{W}_{i,j} \sim \text{Uniform}\left(-\sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}, \sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}\right)$$

$$\mathbf{W}_{i,j} \sim \mathcal{N}\left(0, \sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}\right)$$

He Initialization

Similar to Xavier Initialization, but for **different activation functions**

Activation function	Uniform distribution $[-r, r]$	Normal distribution
Logistic	$r = \sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = \sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$
Hyperbolic tangent	$r = 4\sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = 4\sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$
ReLU (and its variants)	$r = \sqrt{2}\sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = \sqrt{2}\sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$


```
keras.layers.Dense(10, activation="relu", kernel_initializer="he_normal")
```

2. Activation





Activation Function Choices

- Poor choice of activation function can lead to vanishing/exploding gradient
- **Nature** chooses to use roughly **sigmoid activation** function in biological neurons, but it turns out that other functions (eg. ReLU) behave **much faster and better**.



ReLU Activation





- Fast to compute, but suffer from **dying**: meaning not outputting anything but 0
- To solve this, you can use a variant of ReLU called leaky ReLU

ELU (Exponential Linear Unit)

Outperform ReLU on faster convergence and accuracy (Clevert et al 2015)

$$\text{ELU}_{\alpha}(z) = \begin{cases} \alpha(\exp(z) - 1) & \text{if } z < 0 \\ z & \text{if } z \geq 0 \end{cases}$$



Why ELU is better than ReLU?

Dying neuron? non-zero gradient? smoothness?

```
layer = keras.layers.Dense(10, activation="selu",
```


3. Normalization



Batch Normalization

- He Initialization reduces vanishing/exploding gradient problems at the **beginning** of training, but does not guarantee they won't come back **during** training
- Sergey Ioffe and Christian Szegedy (2015) address this by a technique called Batch Normalization (BN)
- This adds an operation **before** activation function: simply **zero-centering** and **normalizing** the inputs, then **scaling** and **shifting** the results → optimal scale.

Batch Normalization

Learn 4 parameters: **scale**, **offset**, **mean**, and **standard deviation**

$$\mu_B = \frac{1}{m_B} \sum_{i=1}^{m_B} \mathbf{x}^{(i)}$$

$$\sigma_B^2 = \frac{1}{m_B} \sum_{i=1}^{m_B} (\mathbf{x}^{(i)} - \mu_B)^2$$

$$\hat{\mathbf{x}}^{(i)} = \frac{\mathbf{x}^{(i)} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$$\mathbf{z}^{(i)} = \gamma \otimes \hat{\mathbf{x}}^{(i)} + \beta$$

- μ_B is the empirical mean, evaluated over the whole mini-batch B .
- σ_B is the empirical standard deviation, also evaluated over the whole mini-batch.
- m_B is the number of instances in the mini-batch.
- $\mathbf{x}^{(i)}$ is the zero-centered and normalized input.
- γ is the scaling parameter for the layer.
- β is the shifting parameter (offset) for the layer.
- ϵ is a tiny number to avoid division by zero (typically 10^{-3}). This is called a *smoothing term*.
- $\mathbf{z}^{(i)}$ is the output of the BN operation: it is a scaled and shifted version of the inputs.

Batch Normalization Implementation

```
model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28,
    keras.layers.BatchNormalization(),
    keras.layers.Dense(300, activation="el
    keras.layers.BatchNormalization(),
    keras.layers.Dense(100, activation="el
    keras.layers.BatchNormalization(),
    keras.layers.Dense(10, activation="sof
])
```

```
>>> model.summary()
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
flatten_3 (Flatten)	(None, 784)	0
batch_normalization_v2 (Batch Normalization)	(None, 784)	3136
dense_50 (Dense)	(None, 300)	235500
batch_normalization_v2_1 (Batch Normalization)	(None, 300)	1200
dense_51 (Dense)	(None, 100)	30100
batch_normalization_v2_2 (Batch Normalization)	(None, 100)	400
dense_52 (Dense)	(None, 10)	1010
Total params: 271,346		
Trainable params: 268,978		
Non-trainable params: 2,368		

Gradient Clipping

- Use as an alternative to Batch Normalization
- **Quick and dirty:** simply clip the gradients during backpropagation so that they can never exceed some threshold.
- Often use in Recurrent Neural Nets (RNNs)

```
optimizer = keras.optimizers.SGD(clipvalue=1.0)
model.compile(loss="mse", optimizer=optimizer)
```

Recap: Optimization for Deep Neural Nets

1. **Initialization:** how to initialize the weights so that they do not saturate?
2. **Activation:** how to solve the vanishing gradient problem?
3. **Normalization:** how to get the model to learn the optimal scale?
4. **Optimizers:** when gradient descent was too slow or not good enough?
5. **Adaptive Learning Rate:** what if convergence is too slow or sub-optimal?
6. **Second-Order Training Methods:** can we make use of second derivatives?

