

Sequence Modeling

Recurrent Neural Networks and LSTM

N. Rich Nguyen, PhD
SYS 6016

Sequential Data

Text: “I am comfortably watching this lecture at home”


Audio:



Stock prices:

21,917.16 -410.32 (1.84%) ↓

Mar 31, 5:14 PM EDT - Disclaimer



A Sequence Modeling Problem

"I am comfortably watching this lecture at home"

given these words

predict the
next word

We can use a fixed window:

"I am comfortably watching this lecture at home"

given these
two words

predict the
next word

Problem 1: Long-term Dependencies

One-hot encoding tells us what each word is: $[1\ 0\ 0\ 0\ 0\ 0\ 1\ 0\ 0\ 0] \rightarrow$ prediction
lecture at

Using a short window (2 words), we can't model long-term dependencies:

"The coronavirus is originated in China, and We should not call it the _____"

not enough info!

We need information from **the distant past...**

Problem 2: Word Order in the Sequence

What if we use the entire sequence as a bag-of-words vector?

"I am comfortably watching this lecture at"



"bag-of-words"

[0 1 0 0 1 0 1 1 ... 0 0 1 1 0 0 1 0 0 0]

However, bag of words does **not preserve the order** of words in the sequence.



The food was good, not bad at all.

vs.

The food was bad, not good at all.



Problem 3: Parameter Sharing

What if we use a very long fixed window:

"I am comfortably watching this lecture at home"

given these words

predict the next word

[1 0 0 0 0 0 0 0 0 1 ... 0 1 0 0 0 0 0 0 1 0] → prediction
I am ... lecture at

Each of these inputs has separate parameters which *are not shared across the sequence*, so a neural network will likely to learn different parameters to “I am” if it appears at different part of the sequence:

[0 1 0 0 0 ... 1 0 0 0 0 0 0 0 0 1 ...] → prediction?
I am ...

→ needs to share the same parameters across several time steps.

Sequence Modeling Criteria

To model sequence, we will need to:

1. Handle variable-length sequence while tracking long-term dependencies
2. Maintain information about order
3. Share parameters across the sequence


Today, we will use **Recurrent Neural Networks (RNNs)** as an approach to *solve sequence modeling* problems.

Recurrent Neural Networks (RNNs)

Recurrent Neural Network

Want to predict \mathbf{y} at some time step t

We can process a sequence of vector \mathbf{x} by applying a **recurrence** formula at every time step:




$$\mathbf{h}_t = f(\mathbf{h}_{t-1}, \mathbf{x}_t)$$

↑ New state
↑ A function with parameters \mathbf{W}
↑ Old state
↑ Input vector at this time step

Recurrent Neural Network

Note that **the same** function and set of parameter are used at every time step.

The state consists of a single “hidden” vector \mathbf{h}_t :

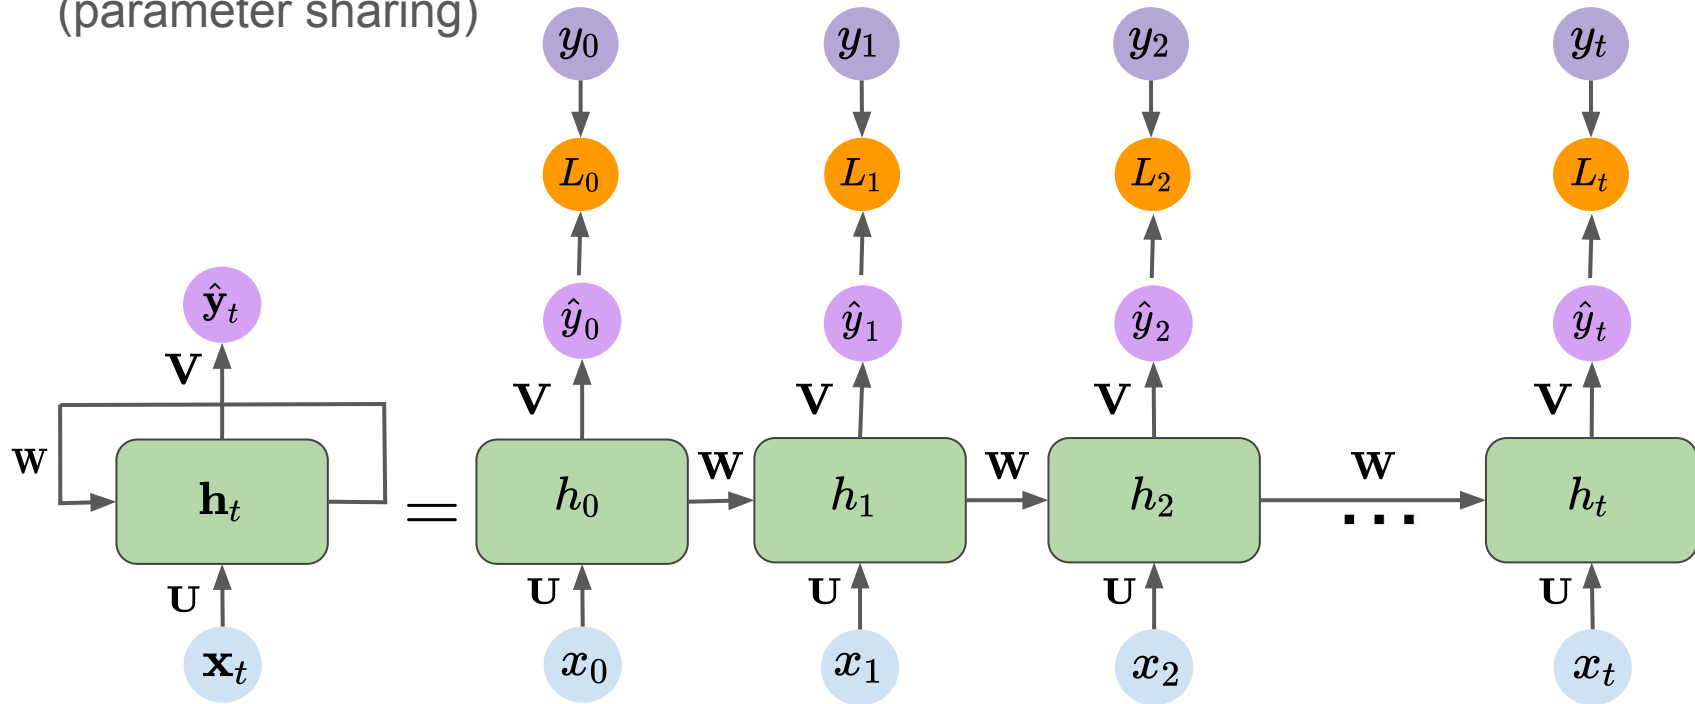


$$\mathbf{h}_t = \tanh(\mathbf{W}\mathbf{h}_{t-1} + \mathbf{U}\mathbf{x}_t)$$

$$\hat{\mathbf{y}}_t = \mathbf{V}\mathbf{h}_t$$

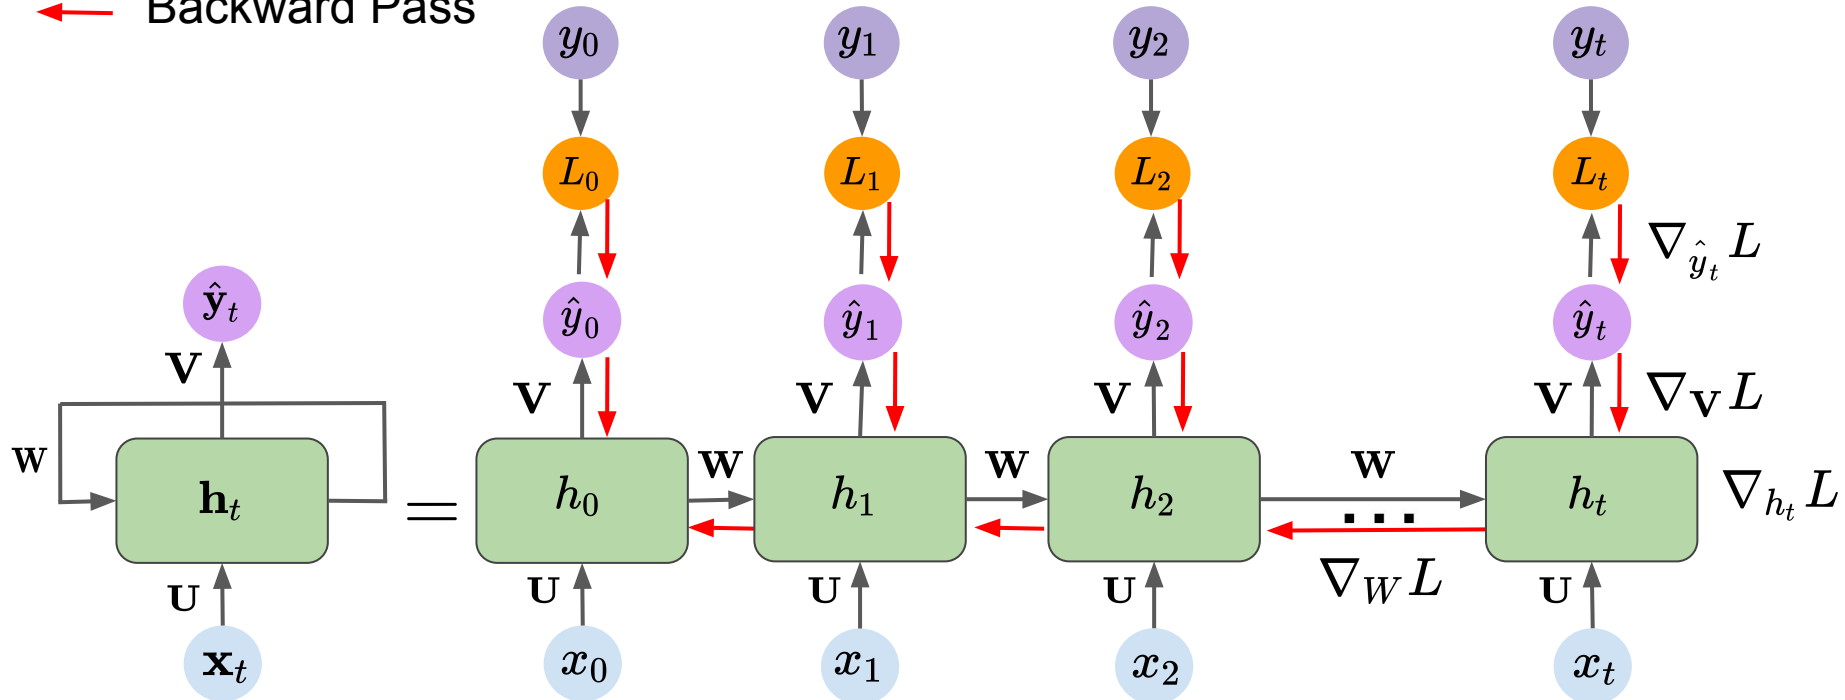
Recurrent Computational Graph

Unfolding the network and **reuse the same weight matrix** at every timestep
(parameter sharing)



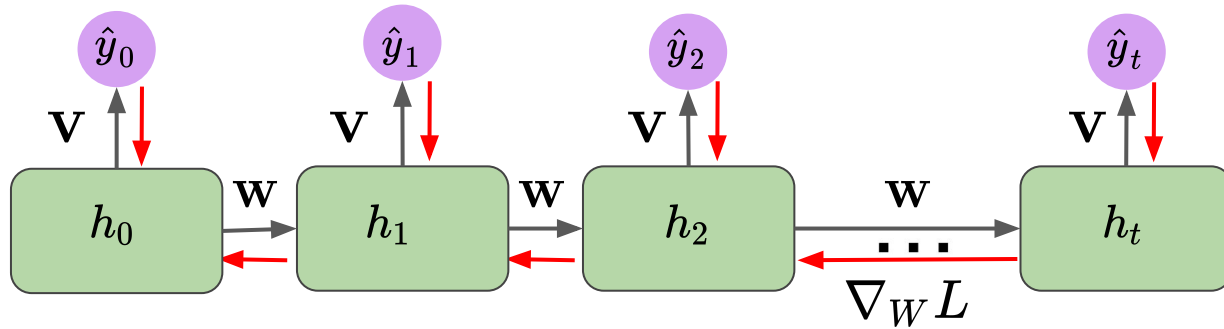
Training RNNs

→ Forward Pass
 ← Backward Pass



Backpropagation Through Time (BPTT)

Training RNN: Vanishing/exploding gradients



Computing the gradient w.r.t h_0 involves many factors of \mathbf{w} with repeated gradient computation which can be problematic:


- Many values > 1 : **exploding** gradients (rarely happen, but damaging) \rightarrow needs to scale big gradients using gradient clipping?
- Many values < 1 : **vanishing** gradients (happen most of the time) \rightarrow only capture short-term dependencies

Problem of Training RNN over long sequence


- A standard RNN can only capture short-term dependencies.
- Training a RNN on long sequences suffer from the **vanishing gradients** problem
- \Rightarrow further back time steps have smaller and smaller gradients
- \Rightarrow this training produces only parameters that can predict short-term dependencies.



“Watching lectures at home is _____”



“COVID-19 is originated from China, and ... , we should not call it the _____”



Control the flow of gradients

Although RNN should theoretically be able to retain information about inputs seen many timesteps before, in practice, such long-term dependencies are very **difficult to learn**²

A solution is to use a more complex recurrent unit with gates to control how the information is passed through


⇒ **Long Short-term Memory (LSTM)**

²See, for example, Yoshua Bengio, Patrice Simard, and Paolo Frasconi, "Learning Long-Term Dependencies with Gradient Descent Is Difficult," *IEEE Transactions on Neural Networks* 5, no. 2 (1994).

Long Short-term Memory (LSTM)

Structure of a Standard RNN

In a standard RNN, repeating modules contain a simple computation node

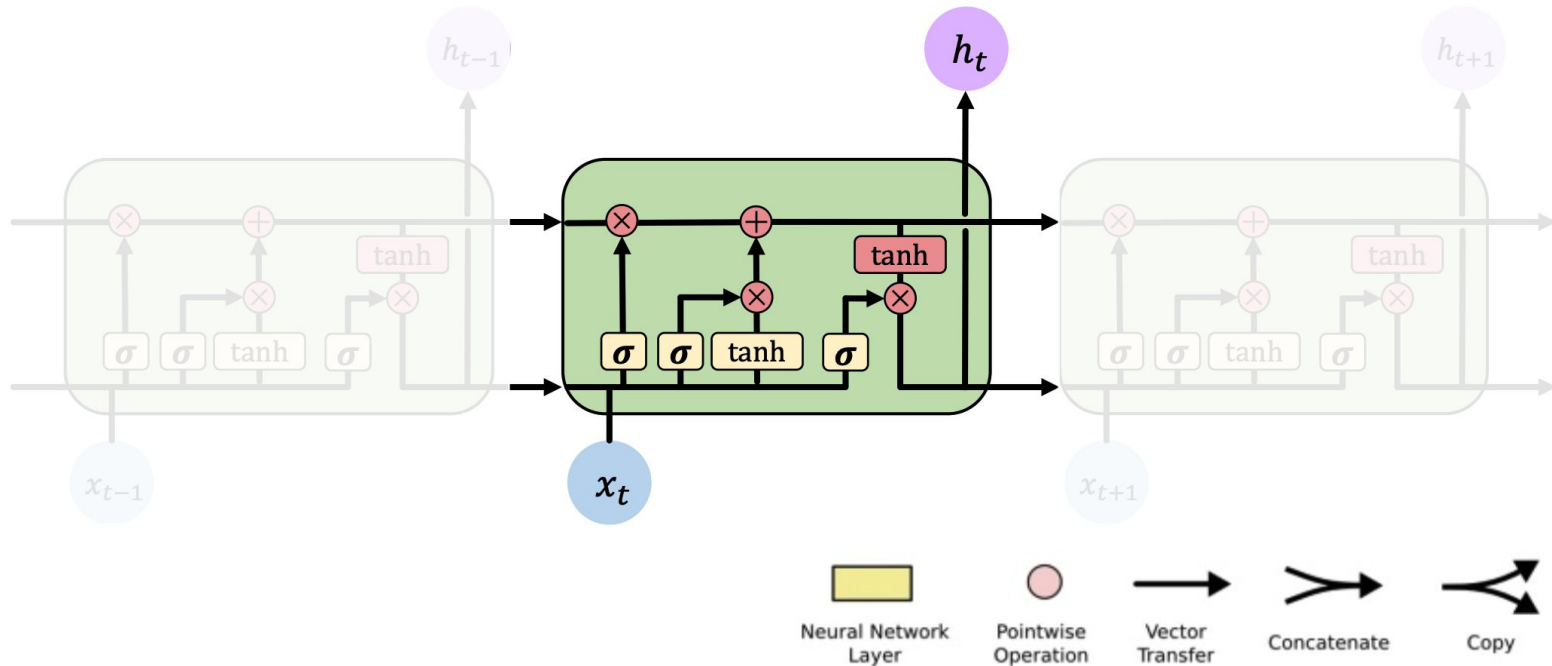


$$\mathbf{h}_t = \tanh(\mathbf{W}\mathbf{h}_{t-1} + \mathbf{U}\mathbf{x}_t)$$

Long Short-Term Memory (LSTM)



Proposed by **Hochreiter** and **Schmidhuber** in 1997

LSTM repeating modules contain interacting layers that control information flow




LSTM Structure

- Introduce **self-loops** to produce paths where the gradient can flow for long duration and make the weight of this self-loop **conditioned** on the context.
- LSTM maintains a **cell state** c_t where it's easy for information flow
- Information is added or removed to cell state through **gates**
- Gates let information through via a **sigmoid** neural layer and pointwise multiplication



LSTM Core Idea



An LSTM cell learns to:


1. **Forget** irrelevant parts of the previous state
2. **Store** relevant parts in a long-term cell state
3. **Update** selectively its cell state
4. **Output** its cell state whenever needed

Example: Bob and Alice are having lunch. Bob likes apples. Alice likes oranges.

She is eating an orange.

1. Forget Gate: Forget irrelevant information

Bob and Alice are having lunch. Bob likes apples. Alice likes oranges. She is eating an orange.




$$\mathbf{f}_t = \sigma(\mathbf{W}_f(\mathbf{h}_{t-1}, \mathbf{x}_t))$$

- Use previous cell output and input
- **Sigmoid** layer: values either 0 or 1
→ “completely forget” or “completely keep”

Example: Forget the gender pronoun of previous subject (Bob)

2. Store Gate: Identify new info to be stored




$$\mathbf{i}_t = \sigma(\mathbf{W}_i(\mathbf{h}_{t-1}, \mathbf{x}_t))$$

$$\tilde{\mathbf{C}}_t = \tanh(\mathbf{W}_C(\mathbf{h}_{t-1}, \mathbf{x}_t))$$

- **Sigmoid:** to decide what values to update
- **Tanh:** to generate new vector of values to be added to the cell state.

Example: Add gender of new subject (Alice) to replace that of old subject (Bob).

3. Update Gate: Update cell state




$$\mathbf{C}_t = \mathbf{f}_t * \mathbf{C}_{t-1} + \mathbf{i}_t * \tilde{\mathbf{C}}_t$$

- Apply **forget** operation to previous internal cell state
- Then, **add** new values, scaled by how much we decided to update

Example: *Actually* drop old info and add **new info** about subject's **gender**.

4. Output Gate: Output filtered cell state



$$\mathbf{o}_t = \sigma(\mathbf{W}_o(\mathbf{h}_{t-1}, \mathbf{x}_t))$$


$$\mathbf{h}_t = \mathbf{o}_t * \tanh(\mathbf{C}_t)$$

- **Sigmoid:** to decide what parts of cell state to output
- **Tanh:** squashes values between -1 and 1
- **Multiplication:** output filtered version of cell state

Example: Having seen a subject (Alice), may output info relating to a pronoun (**She**).

LSTM Gradient Flow

Uninterrupted gradient flow!



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

Summary


- **RNNs** is used to solving sequence modeling problem
- **LSTM** maintains a separate cell state and uses gates to control the information flow: forget, store, update, and output gates.
- Backpropagation can be done with **uninterrupted** gradient flow
- Better/simpler architectures are a **hot** topic of research → the rise of **BERT** in 2018 (will be discussed in the next video)

Bonus Slides


Toy Example: A Moody Chef



Apple pie



Burger



Chicken

Moods based on the weather

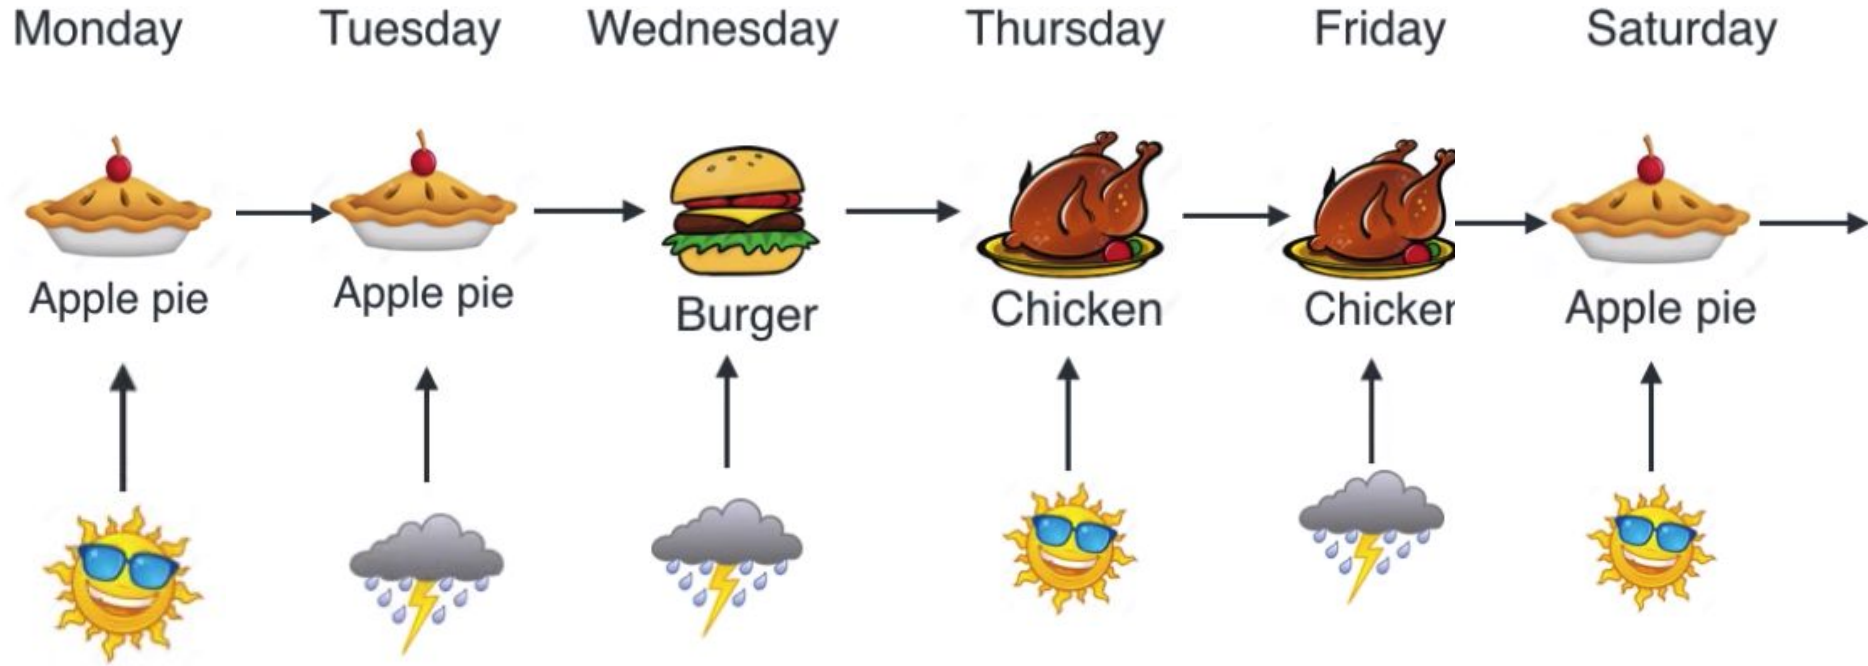


Sunny
Same as yesterday




Rain
Next dish

Cooking Schedule

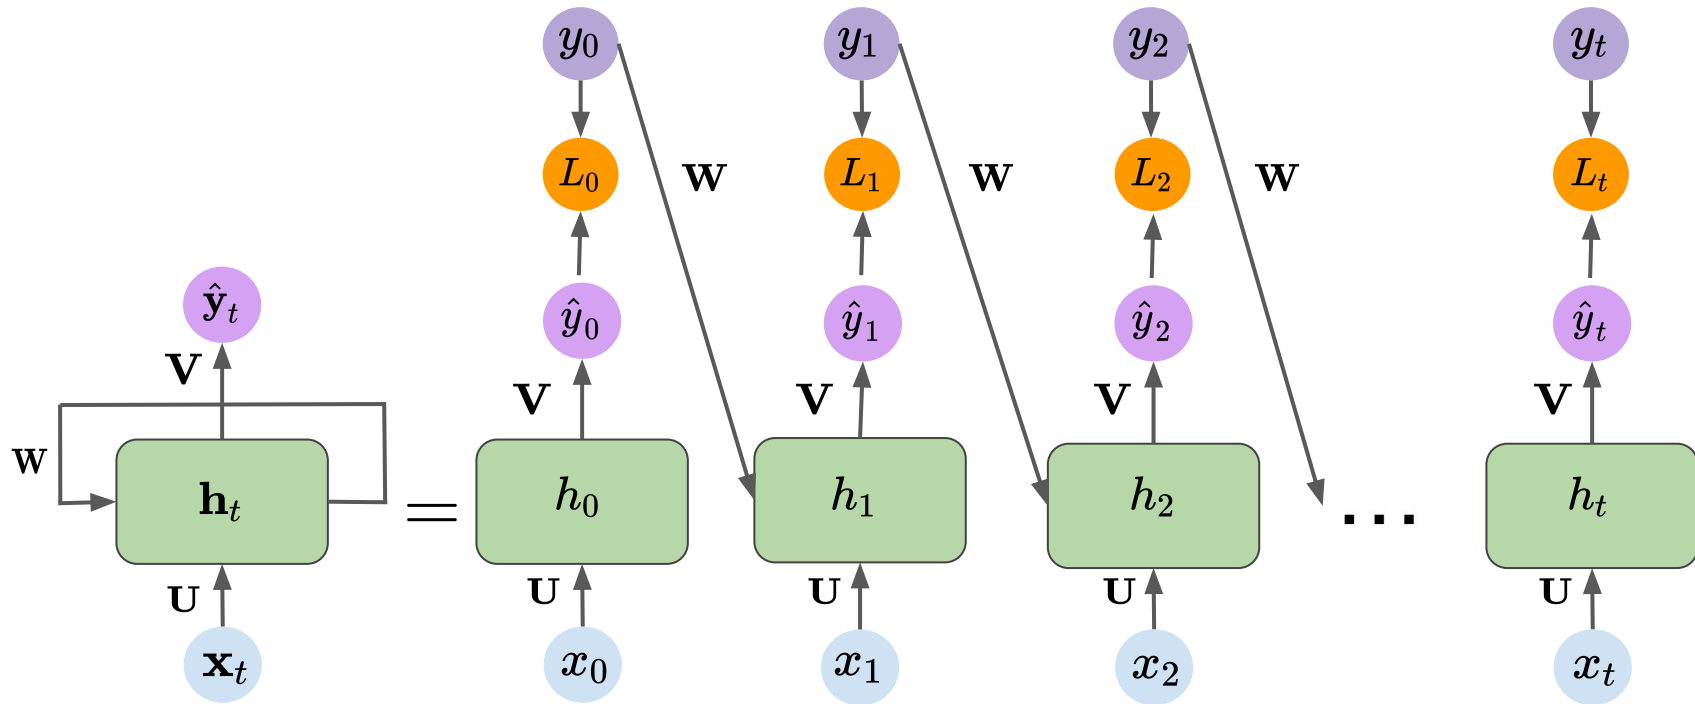


Recurrent Neural Network

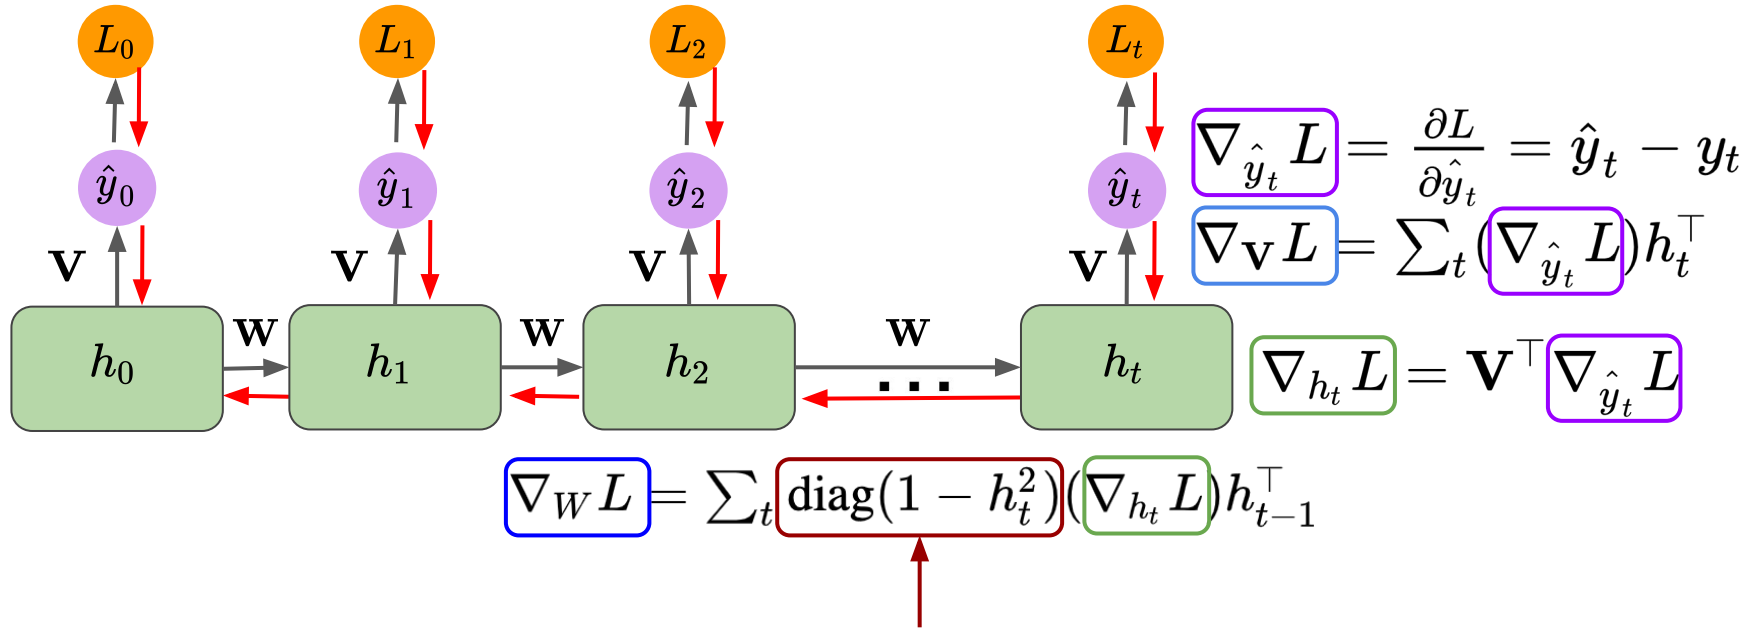


Teacher Forcing Networks

RNN has recurrent connections from **outputs** leading back into **hidden states**



Exploding/vanishing gradients



the Jacobian of the tanh associated
with the hidden unit at time t