

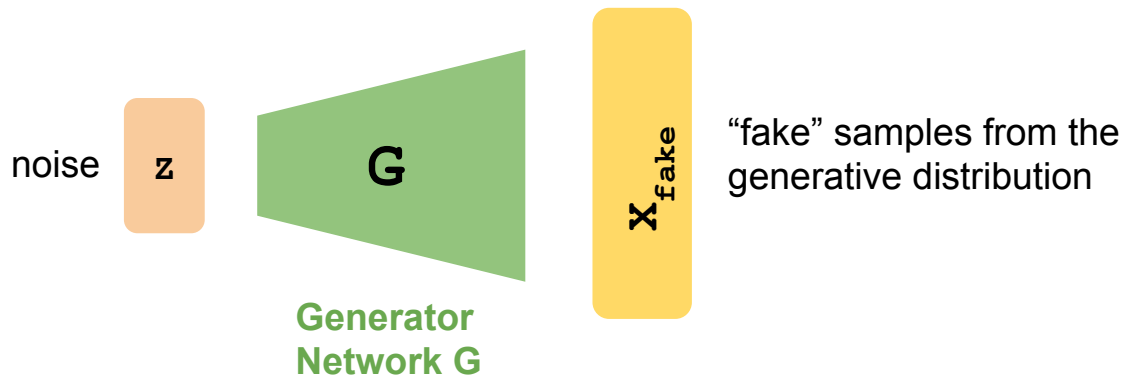
Generative Learning

Advances in Generative Adversarial Networks

N. Rich Nguyen, PhD
SYS 6016

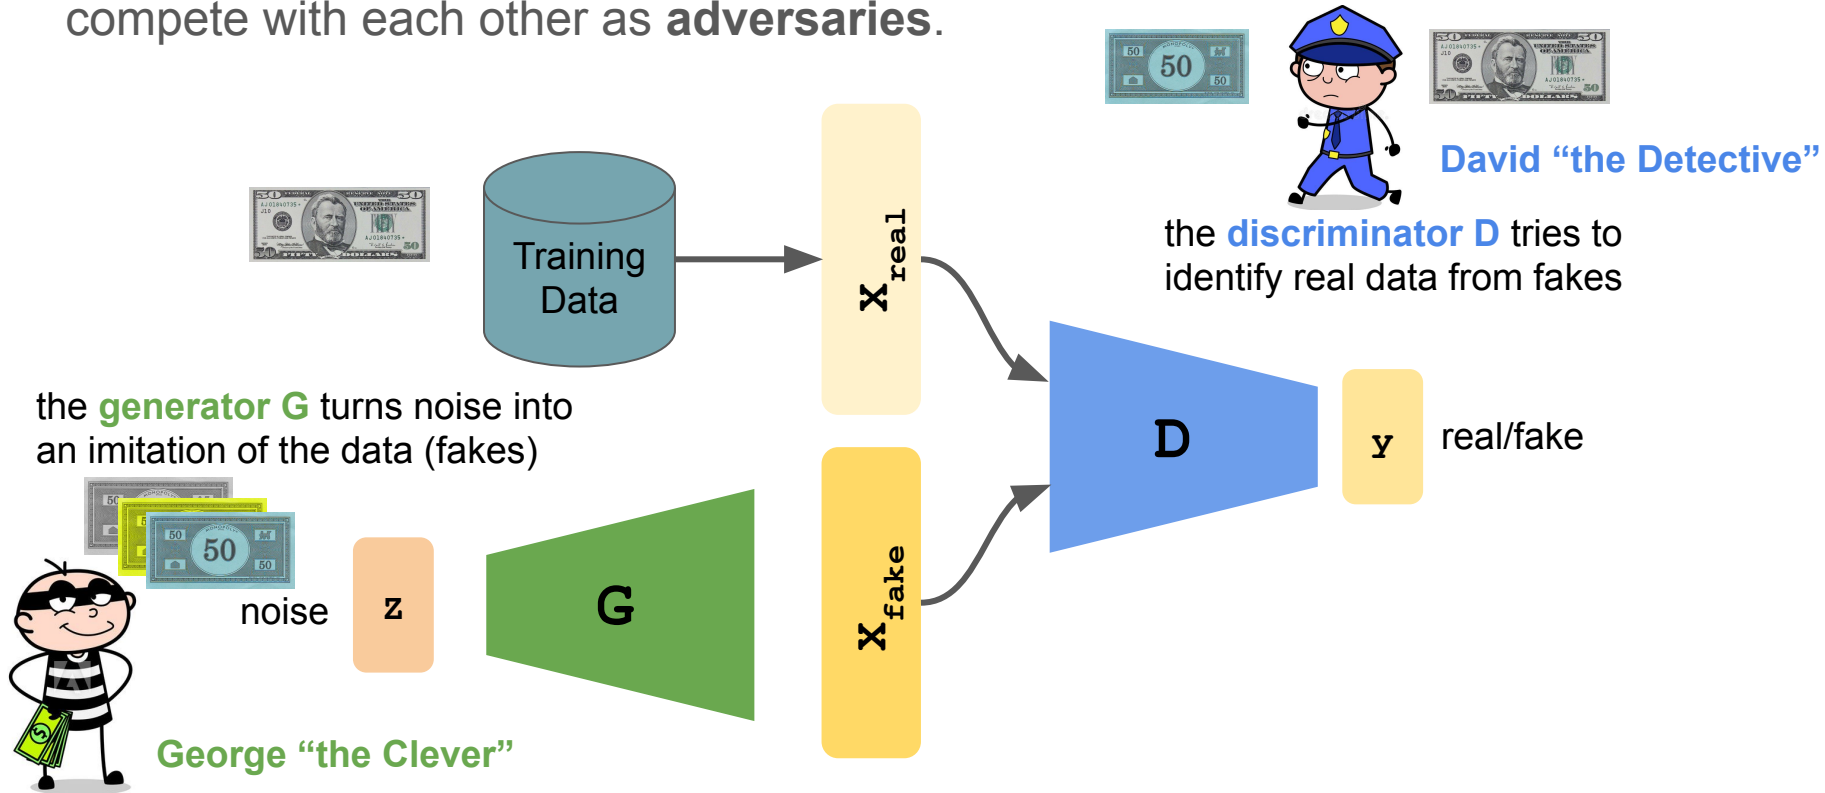
Density distribution: Building or Sampling?

- **Problem:** we tried to generate samples from a complex density distribution, but it's hard to learn such a distribution directly!
- **Solution:** instead of attempt to model the density distribution, just generate new instances by sampling from something simple (like a noise distribution), and then learn a transformation to a desired distribution.



Generative Adversarial Networks (GANs)

GANs are a way to make a generative model by having two neural networks compete with each other as **adversaries**.



Designing a simple GAN for Fashion MNIST

```
codings_size = 30

generator = keras.models.Sequential([
    keras.layers.Dense(100, activation="selu", input_shape=[codings_size]),
    keras.layers.Dense(150, activation="selu"),
    keras.layers.Dense(28 * 28, activation="sigmoid"),
    keras.layers.Reshape([28, 28])
])

discriminator = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    keras.layers.Dense(150, activation="selu"),
    keras.layers.Dense(100, activation="selu"),
    keras.layers.Dense(1, activation="sigmoid")
])

gan = keras.models.Sequential([generator, discriminator])
```

Training GANs



- **Generator G (George)** tries to create fake data to trick **discriminator D (David)**
- **David** tries to identify real data from fakes created by **George**
- **George** and **David** are trained jointly by the below objective function (called a **minimax game**):

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

- **David** wants to **maximize** the objective s.t. **D (x)** close to 1, and **D (G (z))** close to 0.
- **George** wants to **minimize** the same objective s.t. **D (G (z))** close to 1.
- Convergence: **D (x) = D (G (z)) = 0.5** (**David** cannot tell real/fake anymore)

Intuition behind GANs


GANs are trained by simultaneously updating **David** so that he discriminates between real data and fake data which is generated by **George**.

----- Discriminative Distribution (**David**)




..... Data Distribution p_{data}

———— Generative Distribution p_g (**George**)

————→ The mapping of $\mathbf{x} = p_g(\mathbf{z})$



z is sampled uniformly
 x is generated by p_g



Implementing the training of GAN

```
batch_size = 32
dataset = tf.data.Dataset.from_tensor_slices(X_train).shuffle(1000)
dataset = dataset.batch(batch_size, drop_remainder=True).prefetch(1)
```


```
discriminator.compile(loss="binary_crossentropy", optimizer="rmsprop")
discriminator.trainable = False
gan.compile(loss="binary_crossentropy", optimizer="rmsprop")
```

```
def train_gan(gan, dataset, batch_size, codings_size, n_epochs=50):
    generator, discriminator = gan.layers
    for epoch in range(n_epochs):
        for X_batch in dataset:
            # phase 1 - training the discriminator
            noise = tf.random.normal(shape=[batch_size, codings_size])
            generated_images = generator(noise)
            X_fake_and_real = tf.concat([generated_images, X_batch], axis=0)
            y1 = tf.constant([[0.]] * batch_size + [[1.]] * batch_size)
            discriminator.trainable = True
            discriminator.train_on_batch(X_fake_and_real, y1)
            # phase 2 - training the generator
            noise = tf.random.normal(shape=[batch_size, codings_size])
            y2 = tf.constant([[1.]] * batch_size)
            discriminator.trainable = False
            gan.train_on_batch(noise, y2)
```

```
train_gan(gan, dataset, batch_size, codings_size)
```

Generating new data with trained GANs

- After training, use **George** to create new data that's never been seen before.
- If it's trained *effectively*, the generated data will be **indistinguishable** to real data (to both **David** and humans in general)




Images generated after one training epoch



Unfortunately, the images never really get much better than this. It turns out that training a GAN is very **challenging**!

The difficulties of training GANs

- During training, both **George** and **David** constantly try to outsmart each other, in a **zero-sum game**!
- As training advances, the game may end up in a state that game theorists call a **Nash Equilibrium** named after John Nash (played by Russell Crowe in *Beautiful Mind*). Either **George** nor **David** would be better off changing his strategy assuming another player do not change his.
- However, it's not that simple: **nothing guarantees** Nash Equilibrium can ever be reached!




The difficulties of training GANs (cont)

- **Mode Collapse:** this happens when the **George**'s output gradually becomes less diverse. For example, *suppose **George** has successfully fool **David** with image of shoes, and this will encourage him to produce even more images of shoes and forgets to produce anything else.*
- **Hyperparameter:** Because **George** and **David** are constantly compete, their parameters may end up *oscillating* and become *unstable*. GANs are very sensitive to hyperparameters → lots of effort fine tuning them.
- **Experience Replay:** storing images produced by **George** at each iteration in a *replay buffer*, and the training of **David** use images drawn from this buffer instead of produced by the current state of **George**.
- The *dynamics* of **George** and **David** are still not well understood.





Original GANs (2014 by Goodfellow*)

MNIST



Nearest training sample to the generated image on its left.
TFD

CIFAR-10
(Dense Model)



CIFAR-10
(Conv D and Deconv G)

*Goodfellow, Ian, et al. "Generative adversarial nets." *Advances in neural information processing systems*. 2014.

Recent Advances in GANs

GAN quality has progressed rapidly (adopted from a tweet by Ian Goodfellow)




2014



2015



2016



2017



2018

- DCGANs: Deep Convolutional GANs (2015)
- CycleGAN: domain transformation (2017)
- Progressive Growing GANs (2018)
- StyleGAN: style-based generator (2019)

Deep Convolutional GANs (DCGANs*)


Making some guidelines for building a stable convolutional GANs:

- Replace pooling layers with strided convolutions in **David** and transposed convolutions in **George**.
- Use Batch Normalization in both **George** and **David**
- Use ReLU for **George** for all layers except output layer
- Use leaky ReLU in **David** for all layers
- Remove Dense layers

⇒ **Fairly realistic** images, they will get better with more data and deeper networks!

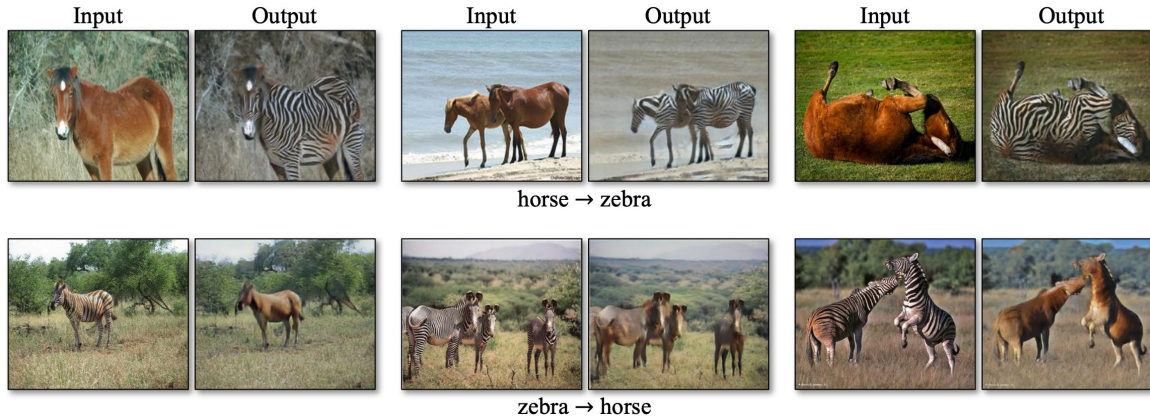



DCGANs: Latent Vector Arithmetics



CycleGAN: domain transformation*


CycleGAN learns transformation across domains with unpaired data



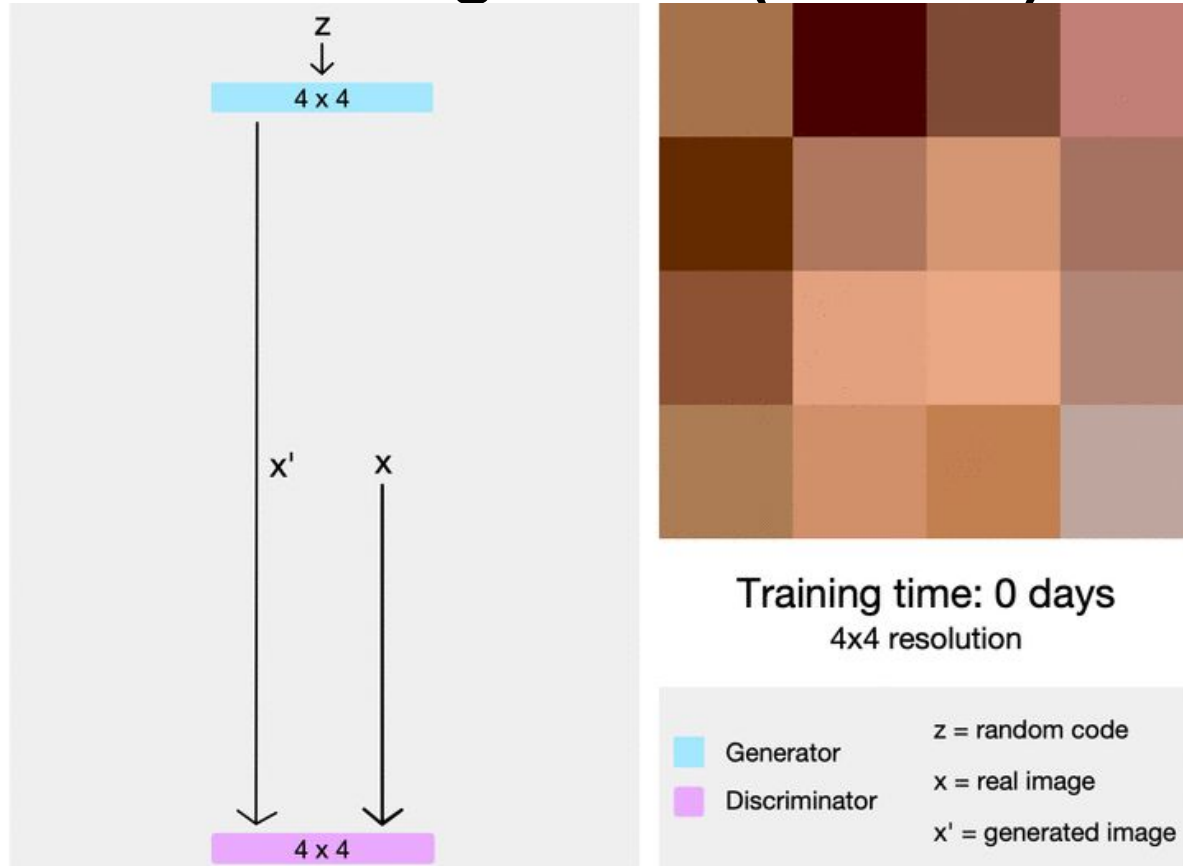
Not only just
adding/removing **stripes**...

Also changes the **grass** to
more yellow and dry
(typically found in zebra's
habitat in African savanna)

CycleGAN demo



Progressive Growing GANs (NVIDIA)*



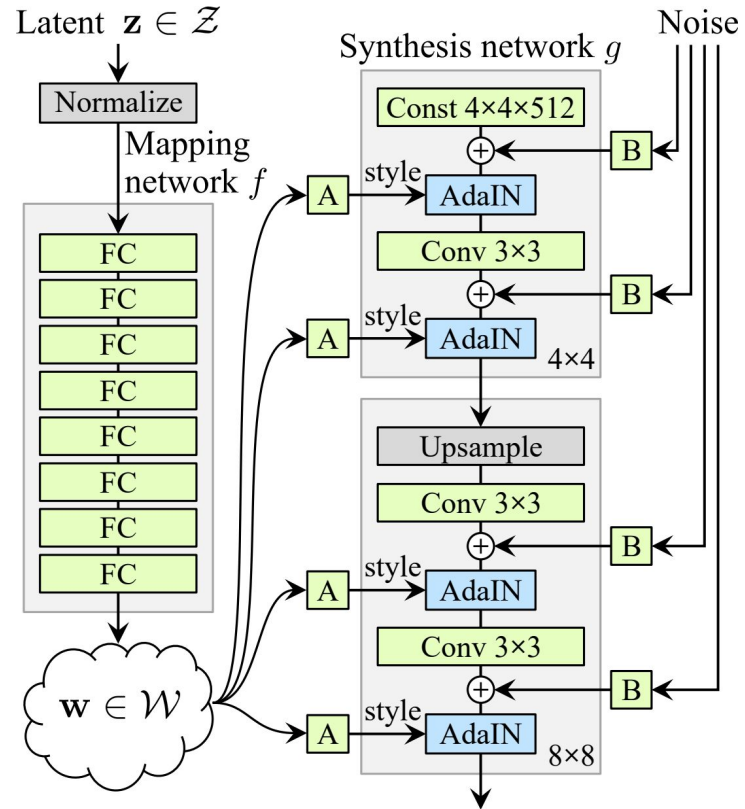
*Karras, Tero, et al. "Progressive Growing of GANs for Improved Quality, Stability, and Variation, ICLR 2018."

Progressive Growing GANs: on CelebA-HQ




StyleGAN: Style-based generator

- Use style transfer techniques in generator **George** to ensure generated images have the same local structure as the training images.
- The discriminator **David** and loss function were not modified.
- **George** now consists of two networks:
 - **Mapping Network**: maps the latent representation z to multiple style vectors
 - **Synthesis Network**: generates images through multiple convolutional and upsampling layers.
- Adding noise independently



StyleGAN: Style-based generator




StyleGAN: Style-based transfer results



Image2StyleGAN


Ryan Gosling \Rightarrow Barack Obama



Oprah Winfrey \Rightarrow Scarlett Johansson

Summary: Generative Adversarial Network

- GANs are “*the most interesting idea in the last ten years in ML*” -- Yann LeCun
- GANs generate *realistic* images that are **indistinguishable** from real images
- The *dynamics* of adversaries **George** and **David** are still not *well* understood
- There are such a variety of GANs out there and still growing
- This hopefully gives you an introduction and perhaps a desire to learn more!



Bonus Content

DCGANs: Implementation

```
codings_size = 100

generator = keras.models.Sequential([
    keras.layers.Dense(7 * 7 * 128, input_shape=[codings_size]),
    keras.layers.Reshape([7, 7, 128]),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2DTranspose(64, kernel_size=5, strides=2, padding="same",
                                   activation="selu"),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2DTranspose(1, kernel_size=5, strides=2, padding="same",
                                   activation="tanh")
])

discriminator = keras.models.Sequential([
    keras.layers.Conv2D(64, kernel_size=5, strides=2, padding="same",
                        activation=keras.layers.LeakyReLU(0.2),
                        input_shape=[28, 28, 1]),
    keras.layers.Dropout(0.4),
    keras.layers.Conv2D(128, kernel_size=5, strides=2, padding="same",
                        activation=keras.layers.LeakyReLU(0.2)),
    keras.layers.Dropout(0.4),
    keras.layers.Flatten(),
    keras.layers.Dense(1, activation="sigmoid")
])

gan = keras.models.Sequential([generator, discriminator])
```