

# Reinforcement Learning

Introduction, Markov Decision Process, and  
Q-Learning

N. Rich Nguyen, PhD  
**SYS 6016**

# Classes of Learning Problems

## Supervised Learning:

**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:** Learn function to map  $x \rightarrow y$

**Example:**



This thing is an apple.

## Unsupervised Learning:

**Data:**  $x$

$x$  is data, no labels!

**Goal:** Learn underlying structure

**Example:**



This thing is like the other thing.

## Reinforcement Learning:

**Data:** state-action pairs

**Goal:** Maximize future rewards over many steps

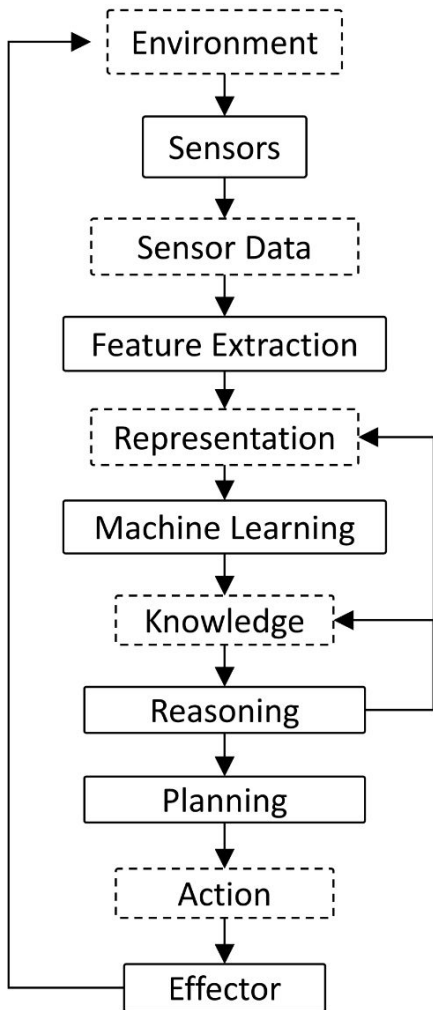
**Example:**



Eat this thing because it will keep you alive.


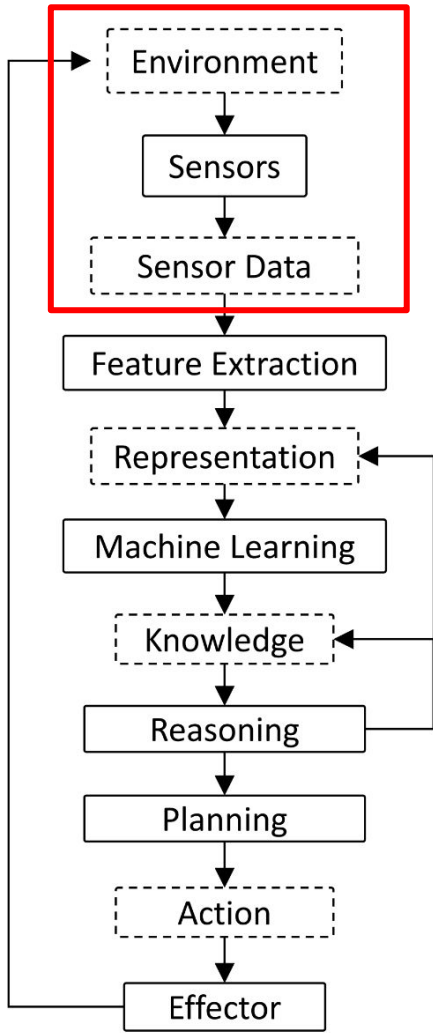


# The Machine Learning Stack!



What can be learned?

# Sensors



Lidar



Camera  
(Visible, Infrared)



Radar



GPS




Stereo Camera



Microphone

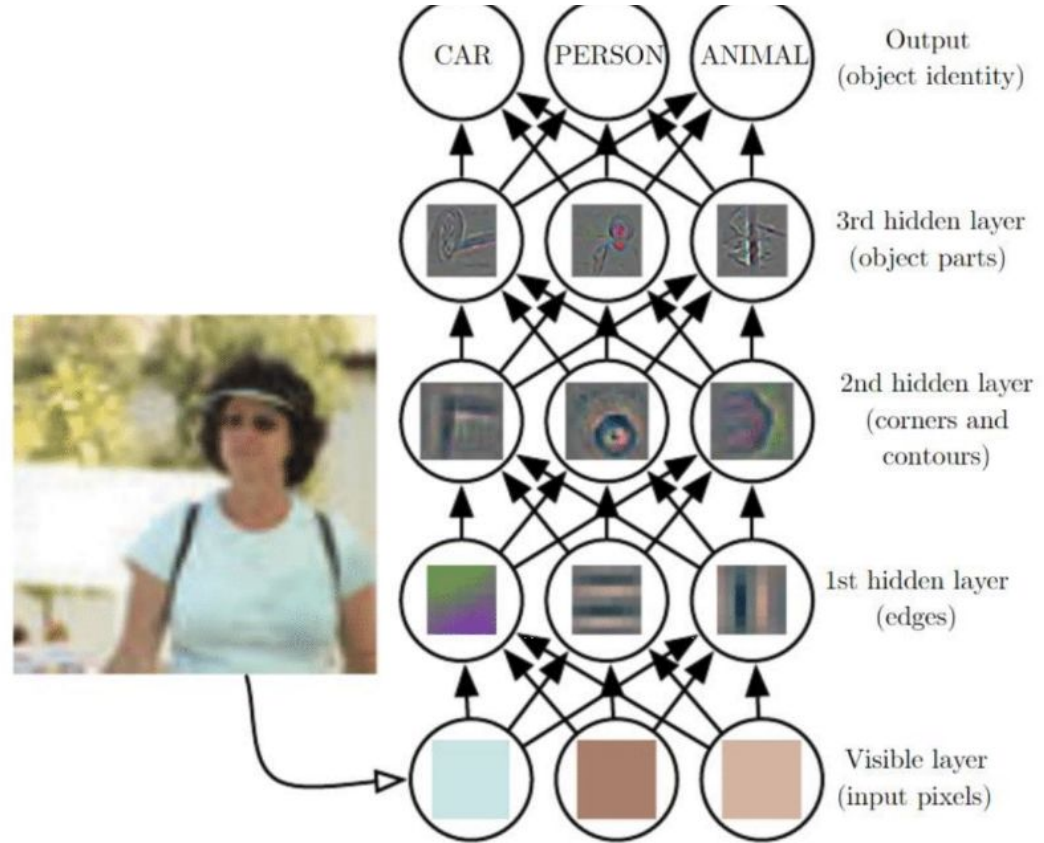
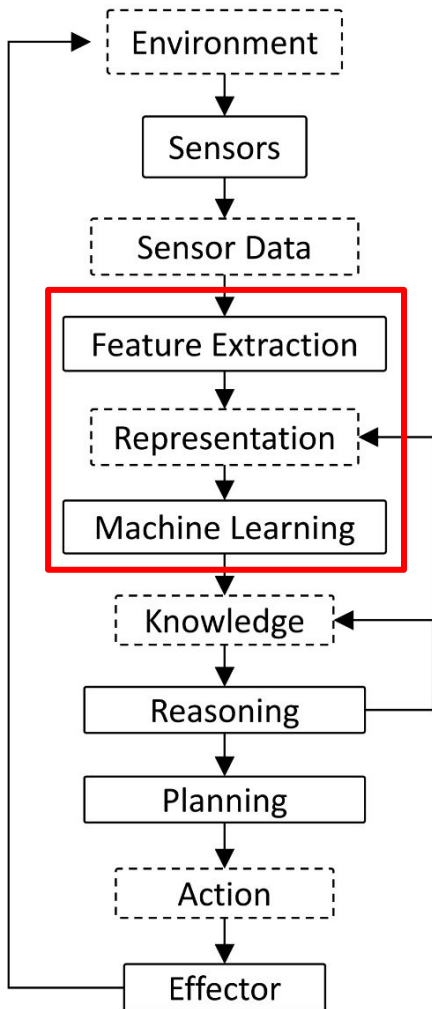


Networking  
(Wired, Wireless)



IMU

# Representations




# Knowledge / Reasoning

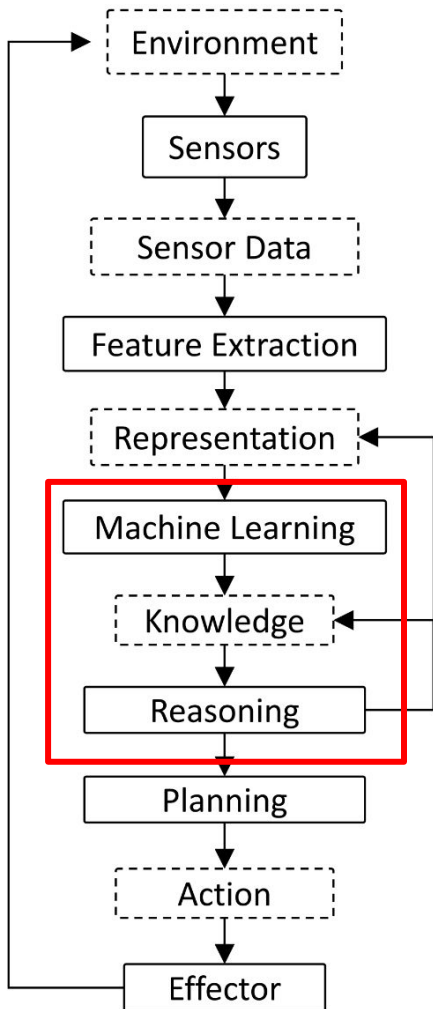
**Image Recognition:**  
If it looks like a duck



**Audio Recognition:**  
Quacks like a duck





**Activity Recognition:**  
Swims like a duck



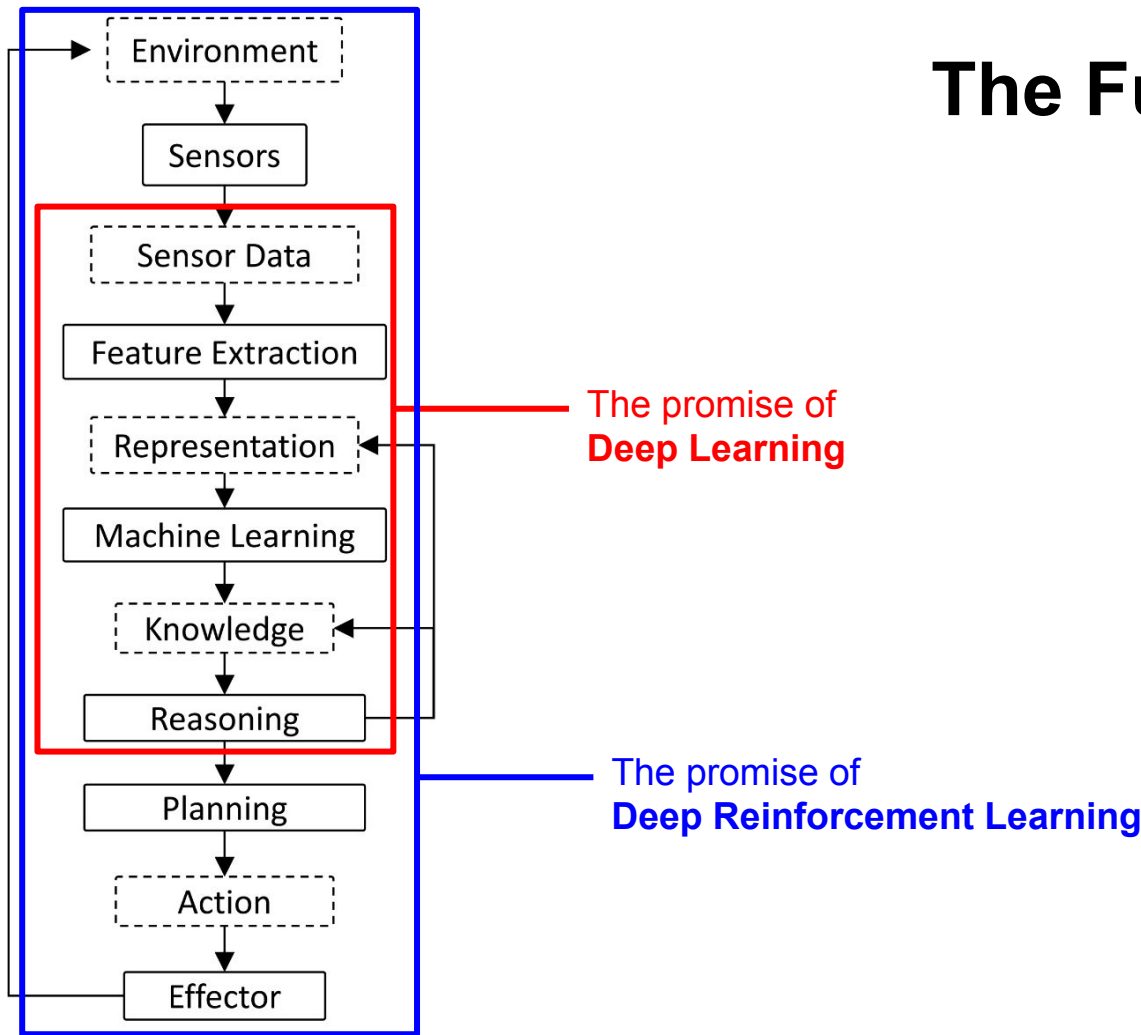


# Actions





# The Full Stack




# Reinforcement Learning

# Reinforcement Learning (RL)

Problems involving an **agent** interacting with an **environment**, which provides numeric reward signals:

At each step, the agent:


- Executes an **action**
- Observes a new **state**
- Receives some **reward**



**Goal:** Learn to take actions from a policy in order to **maximize** some reward


# Example: Cart-Pole Balancing Problem

- **Goal**: Balance the pole of top of a moving cart
- **State**: Pole angle, angular speed, cart position, horizontal velocity
- **Actions**: Horizontal force to the cart
- **Reward**: +1 at each time step if the pole is upright



# Example: Grasping Objects

- **Goal**: Pick an object with different shape
- **State**: Raw pixels from camera
- **Actions**: Move arm, grasp, and lift
- **Reward**: positive when a pickup is successful



# Example: Doom

- **Goal**: eliminate all opponents
- **State**: raw game pixel of the game
- **Actions**: Shoot, Dodge, Run, Walk, Left, Right, No action, ect.
- **Reward**: positive when eliminating an opponent, negative when the agent is eliminated



# Example: MIT DeepTraffic

- **Goal**: train an RL agent that can successfully navigate through traffic
- **State**: as a grid, where each cell will be the speed of the vehicle inside it
- **Actions**: Accelerate, Break, Left, Right, No action
- **Reward**: positive when high speed is maintained.





# RL in Humans





Humans appear to learn to walk through “very few examples” of trial and error.

“**How** we learn how to walk” is an open question...some possible answers:

- **Hardware:** 230M years of bipedal movement data
- **Imitation Learning:** Observation of other humans walking
- **Algorithms:** probably better than backprop and SGD

# RL Real-world Applications

- a. Robotic Explorer
- b. Ms. Pac-man
- c. Go player
- d. Thermostat
- e. Automated Trader



# Markov Decision Process

# Markov Decision Process (MDP)

Markov Decision Process is the mathematical formulation for RL problem

**Markov property:** current state completely characterizes the state of the world.

$$(\mathcal{S}, \mathcal{A}, \mathcal{R}, \mathbb{P}, \gamma)$$


Set of possible states

Set of possible actions

Reward distribution given (state, action)

Transition probability distribution of next state

Discount factor

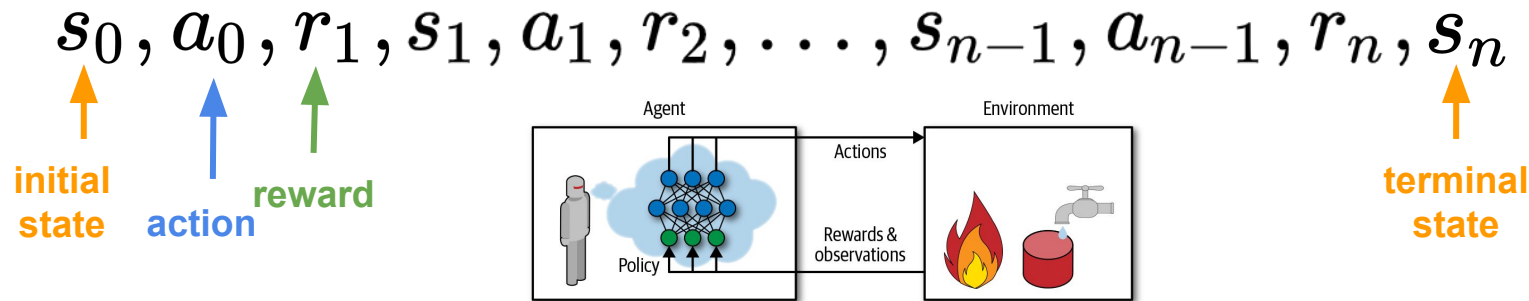


# Markov Decision Process

At time step  $t=0$ , environment samples initial state  $s_0 \sim \mathbb{P}(s_0)$ , then repeat:

- Agent selects **action**  $a_t$
- A **reward** is sampled from the environment  $r_t \sim \mathcal{R}(\cdot | s_t, a_t)$
- Next **state** is sampled from the environment  $s_{t+1} \sim \mathbb{P}(\cdot | s_t, a_t)$
- Agent receives the **reward**  $r_t$  and **next state**  $s_{t+1}$

**Objective:** find a strategy that maximizes **future reward**



# Maximize Future Reward

Future Reward:  $R_t = r_t + r_{t+1} + r_{t+2} + \dots + r_n$

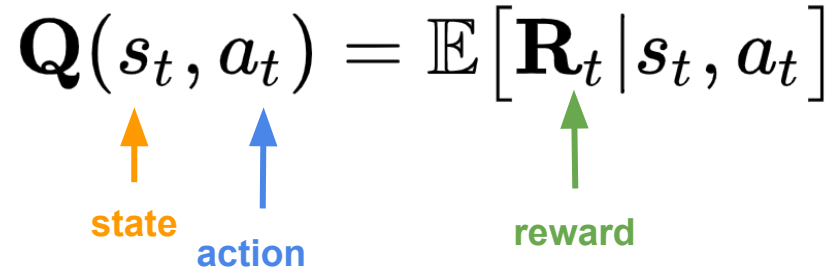
**Discounted Future Reward:**  $R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^{n-t} r_n$   
 $0 < \gamma < 1$

A good strategy for **an agent** would be to always choose **an action** that **maximizes** the **discounted future reward**

Why?

- Uncertainty due the environment, partial observability
- Real life example: Either live it up today, or save \$ for tomorrow?

# Taking actions given a Q-function

$$Q(s_t, a_t) = \mathbb{E}[\mathbf{R}_t | s_t, a_t]$$


The diagram shows the equation  $Q(s_t, a_t) = \mathbb{E}[\mathbf{R}_t | s_t, a_t]$ . Below the equation, three arrows point upwards to specific parts: an orange arrow points to  $s_t$  and is labeled "state" in orange; a blue arrow points to  $a_t$  and is labeled "action" in blue; and a green arrow points to  $\mathbf{R}_t$  and is labeled "reward" in green.

A policy  $\pi$  is a function from **S** to **A** that specifies what action to take in each state

The agent needs a policy  $\pi(s)$  to infer the **best action** to take at a **state s**.

The policy should choose an action that **maximizes the expected future reward (Q-value)**:

$$\pi(s) = \arg \max_a Q(s, a)$$




# Moving in a grid world


**Goal:** reach one of the terminal states (stars) using the least number of actions.

actions = {

1. right 

2. left 

3. up 

4. down 

}


states

★			
			★


Set a negative “reward”  
for each transition  
(e.g.  $r = -1$ )

# Policy to move in a grid world

**Goal:** reach one of the terminal states (stars) using the least number of actions.



Random Policy



Optimal Policy

# 3 Types of Reinforcement Learning

RL agents may be directly or indirectly trying to do one of the following:

1. **Model Learning**: model the environment then use and update it often (difficult because of the stochastic nature of the environment)
2. **Value Learning (Q-Learning)**: learn the state or state-action value, act by choosing best action, and explore if necessary
3. **Policy Learning (Policy Gradients)**: learn the stochastic policy function that maps state to action, act by sampling that policy

Better  
Sample Efficient

Less  
Sample Efficient



Model-based  
(100 time steps)

Off-policy  
Q-learning  
(1 M time steps)

Actor-critic

On-policy  
Policy Gradient  
(10 M time steps)

Evolutionary/  
gradient-free  
(100 M time steps)

# Value Learning / Q-learning

# Q-Learning for Solving the Optimal Policy

Use a **state-action value function**  $Q_{\pi}(s,a)$  as **the expected return** when starting in  $s$ , performing  $a$ , and following  $\pi$

**Q-Learning:** Use any policy to estimate  $Q$  that maximizes future reward:

- $Q$  directly approximates  $Q^*$  (Bellman's equation)
- Independent of the policy, only requirement is keep updating  $(s_t, a_t)$  pair

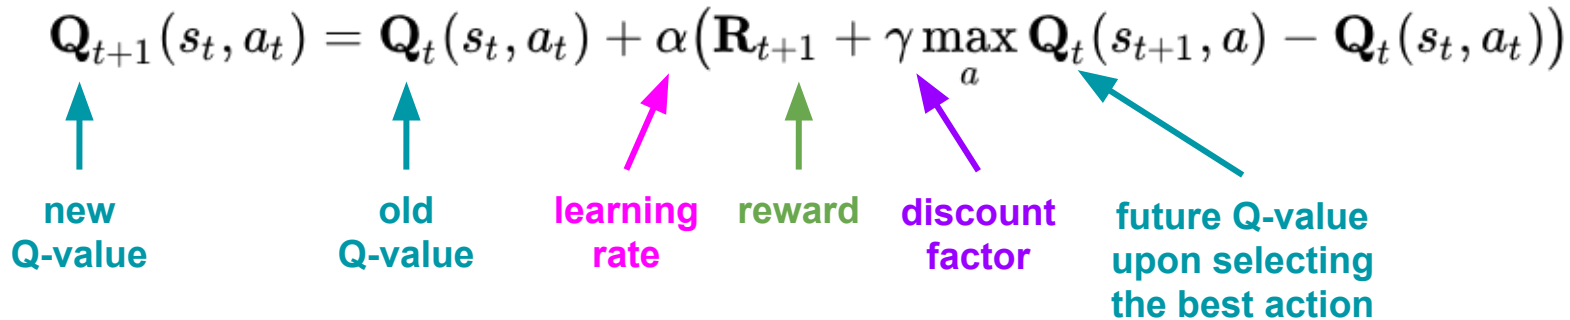
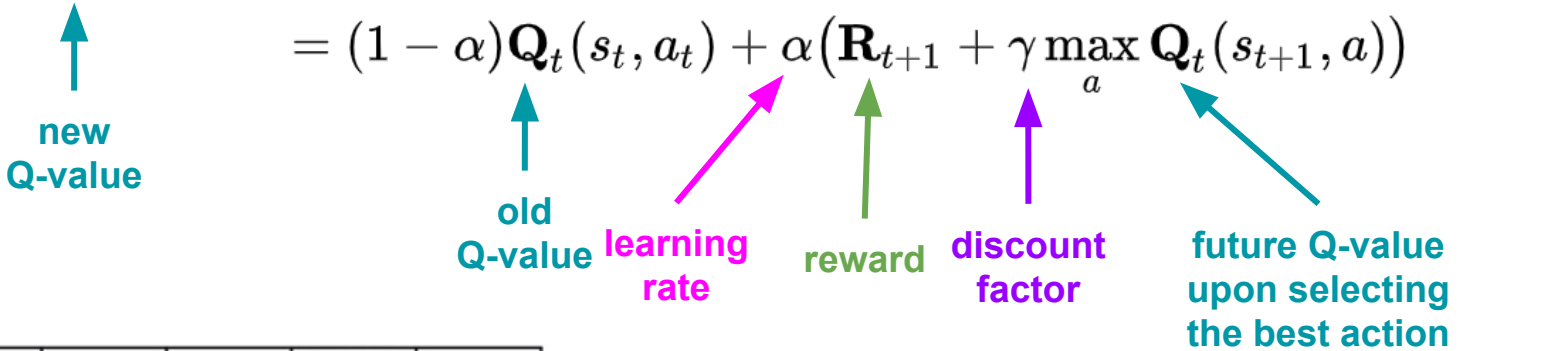
$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha(R_{t+1} + \gamma \max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t))$$


Diagram illustrating the Q-learning update equation with color-coded components:

- new Q-value** (blue arrow pointing to  $Q_{t+1}(s_t, a_t)$ )
- old Q-value** (blue arrow pointing to  $Q_t(s_t, a_t)$ )
- learning rate** (pink arrow pointing to  $\alpha$ )
- reward** (green arrow pointing to  $R_{t+1}$ )
- discount factor** (purple arrow pointing to  $\gamma$ )
- future Q-value upon selecting the best action** (teal arrow pointing to  $\max_a Q_t(s_{t+1}, a)$ )

# Q-Learning: Value Iteration

$$\begin{aligned}
 Q_{t+1}(s_t, a_t) &= Q_t(s_t, a_t) + \alpha(R_{t+1} + \gamma \max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t)) \\
 &= (1 - \alpha)Q_t(s_t, a_t) + \alpha(R_{t+1} + \gamma \max_a Q_t(s_{t+1}, a))
 \end{aligned}$$



	A1	A2	A3	A4
S1	+1	+2	-1	0
S2	+2	0	+1	-2
S3	-1	+1	0	-2
S4	-2	0	+1	+1

```


initialize Q[num_states,num_actions] arbitrarily
observe initial state s
repeat
    select and carry out an action a
    observe reward r and new state s'
    Q[s,a] = Q[s,a] + α(r + γ maxa' Q[s',a'] - Q[s,a])
    s = s'
until terminated
  
```

# Example: Moving in a grid world

Objective: reach one of the terminal states (stars) using the least number of actions.

Find  $Q(s, a)$


$$a = \arg \max_a Q(s, a)$$



actions = {

1. right →
2. left ←
3. up ↑
4. down ↓

}



Optimal Policy




# Exploration vs. Exploitation

Deterministic/greedy policy won't explore all actions:

- Don't know anything about the environment at the beginning
- Need to try all actions to find the optimal one


## $\epsilon$ -greedy policy

- With probability  $(1 - \epsilon)$  of performing the greedy action, otherwise random action
- Slowly move toward greedy policy:  $\epsilon \rightarrow 0$



# Exploration vs. Exploitation Examples

- **Restaurant Selection**
  - **Exploitation:** Go to your favourite restaurant
  - **Exploration:** Try a new restaurant online
- **Banner Ads**
  - **Exploitation:** Show the most successful ads
  - **Exploration:** Show a different ads
- **Oil Drilling**
  - **Exploitation:** Drill at the best known location
  - **Exploration:** Drill at a new location
- **Game Playing**
  - **Exploitation:** Play the move you believe is best
  - **Exploration:** Play an experimental move



# Summary

- Reinforcement Learning (RL) is one of the most exciting fields of ML
- In Markov Decision Process, a good strategy for an agent would be to always choose an action that maximizes the *discounted future reward*
- Q-Learning is to learn the state-action value, act by choosing best action, and explore if necessary

**Next: Deep Q-Networks, Policy Gradients, and Actor Critic Algorithms**

# Acknowledgement

Slides contain some materials and figures reproduced from Lex Fridman (MIT) and Serena Yeung (Stanford) for educational purposes only.

